

# **Hackathon edition#3**

**April 2025**

## **Credit Card Fraud Detection: Model Comparison and Evaluation**

### **Team Details :**

**Manish Tiwari (Leader) (IITMCS\_2406327)**

**Anushka Sarkar (IITMCS\_2406255)**

**Jay Paneliya (IITMCS\_24091607)**

### **Project link**

<https://github.com/JAYpaneliya/credit-card-fraud-detection>

## **Problem Statement :**

In real-world applications, fraud detection remains one of the most critical and complex classification challenges due to highly imbalanced datasets and evolving fraud patterns. The objective of this project was to:

- Train and evaluate multiple machine learning classification models
- Compare their performance on key metrics (Precision, Recall, F1-score, etc.)
- Perform hyperparameter tuning to improve model effectiveness
- Handle class imbalance to ensure realistic fraud detection capabilities
- Build an interactive dashboard for comparison, visualization, and exploration

The challenge emphasized not only model performance but also the ability to present results in a clear and insightful manner using modern tools like Streamlit.

## Dataset Description :

### Source :

The dataset used for this project is the **Credit Card Fraud Detection** dataset made publicly available on Kaggle. It contains real transaction data collected by European cardholders in September 2013. The dataset was downloaded directly from the Kaggle platform.

<https://www.kaggle.com/datasets/mlg-ulb/creditcardfraud/data>



### Dataset Overview

- **Total Records:** 284,807 transactions
- **Fraudulent Transactions:** 492
- **Non-Fraudulent Transactions:** 284,315
- **Imbalance Ratio:** ~0.172% fraud (highly imbalanced)



### Key Features

- Features are result of **PCA anonymization**, named **V1** to **V28**
- Two original features:
  - **Time:** Seconds elapsed between the transaction and the first transaction in the dataset
  - **Amount:** Transaction amount
- **Target Variable:**
  - **Class:** 0 = Non-Fraud, 1 = Fraud

### Preprocessing Steps:

- Scaled **Amount** and **Time** using **StandardScaler**
- **Removed missing values** (none present in this dataset)
- **Applied SMOTE** (Synthetic Minority Oversampling Technique) to balance the dataset before training

## Our Approach

To address the challenge of credit card fraud detection on a highly imbalanced dataset, we followed a structured, multi-step machine learning pipeline:

### Step 1: Data Preprocessing

- **Feature Scaling:** The **Amount** and **Time** features were scaled using **StandardScaler** to normalize their ranges.
- **Missing Values:** Dataset was checked for null or missing values; none were found.
- **Class Imbalance Handling:**
  - The dataset is highly imbalanced (~0.17% fraud cases).
  - We applied **SMOTE (Synthetic Minority Over-sampling Technique)** to the training data to synthetically generate minority class samples and improve model learning.

### Step 2: Train-Test Split

- An **80/20 split** was used to divide the dataset.
  - 80% for training and SMOTE balancing
  - 20% untouched for testing real-world performance

### Step 3: Model Selection

We selected and trained **five popular classification algorithms**:

Model	Description
Logistic Regression	A linear baseline classifier
Random Forest	Esemble method using decision trees
SVM	Kernal based non-linear classifier
Gradient Boosting	Boosted decision trees
XGBoost	High performance boosting algorithm

Each model was trained on the **SMOTE- balanced training set**.

### Evaluation Strategy :

We evaluated all models on the original, imbalanced test set using:

- **Precision** (focus on reducing false positives)
- **Recall** (focus on catching actual frauds)
- **F1-score** (balance between precision & recall)
- **Confusion Matrix**
- **ROC Curve & AUC Score**

This multi-metric evaluation ensured a balanced view of model strengths, especially under imbalanced conditions.

## Hyperparameter Tuning :

We applied **RandomizedSearchCV** on the XGBoost model using:

- F1-score as the optimization metric
- 3-fold cross-validation
- Multiple combinations of learning rate, depth, and estimators

This tuning significantly improved performance and reduced false negatives.

## Implementation Details

Our solution was implemented entirely using Python, with a focus on machine learning model development, evaluation, and visualization. The workflow was developed using **Google Colab** and later integrated into a **Streamlit dashboard** for interactive presentation.

## Technologies and Libraries Used

Tool/Library	Purpose
<b>Pandas</b>	Data loading and manipulation
<b>NumPy</b>	Numerical operations
<b>Scikit-learn</b>	Model training, SMOTE, metrics, tuning
<b>XGBoost</b>	High-performance boosting model
<b>Matplotlib / Seaborn</b>	Visualizations for evaluation
<b>SHAP</b>	Model interpretability (optional)
<b>Streamlit</b>	Interactive dashboard interface
<b>Joblib</b>	Saving and loading trained models



## **Model Training & Saving :**

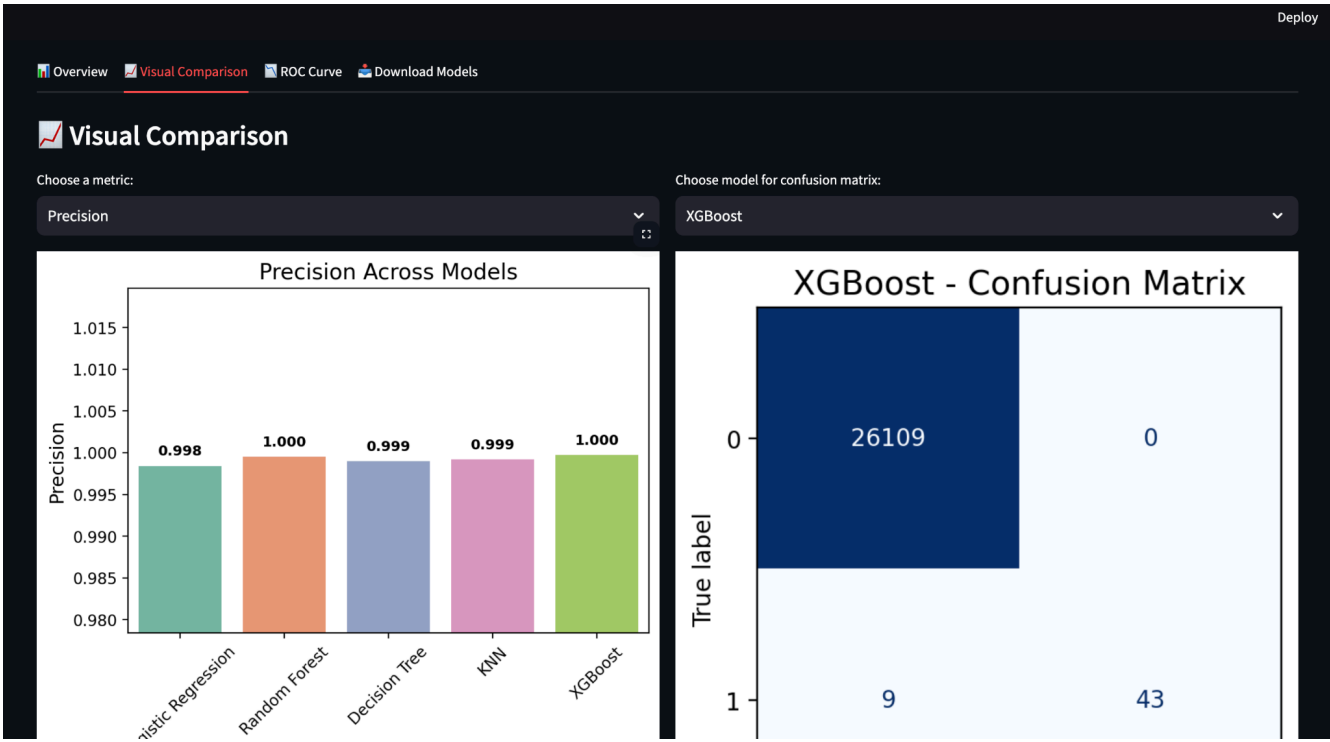
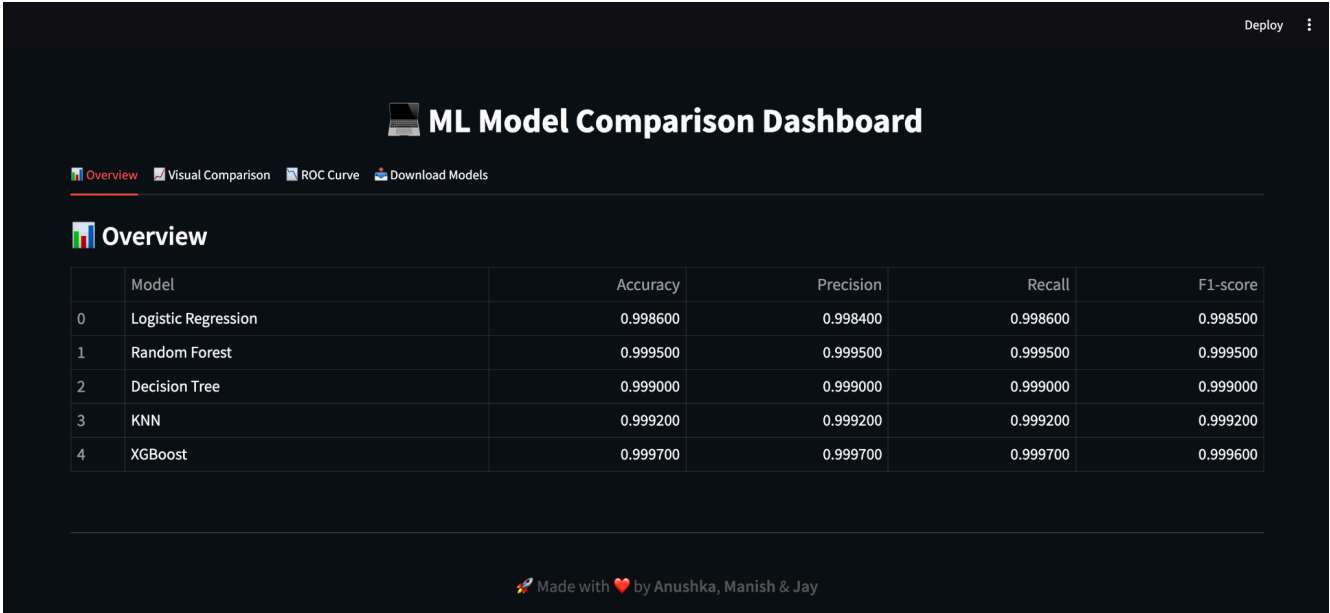
Each model was trained on the SMOTE-balanced dataset and evaluated using the imbalanced test set. Trained models were saved as `.pk1` files using `joblib.dump()` and loaded later into the dashboard.

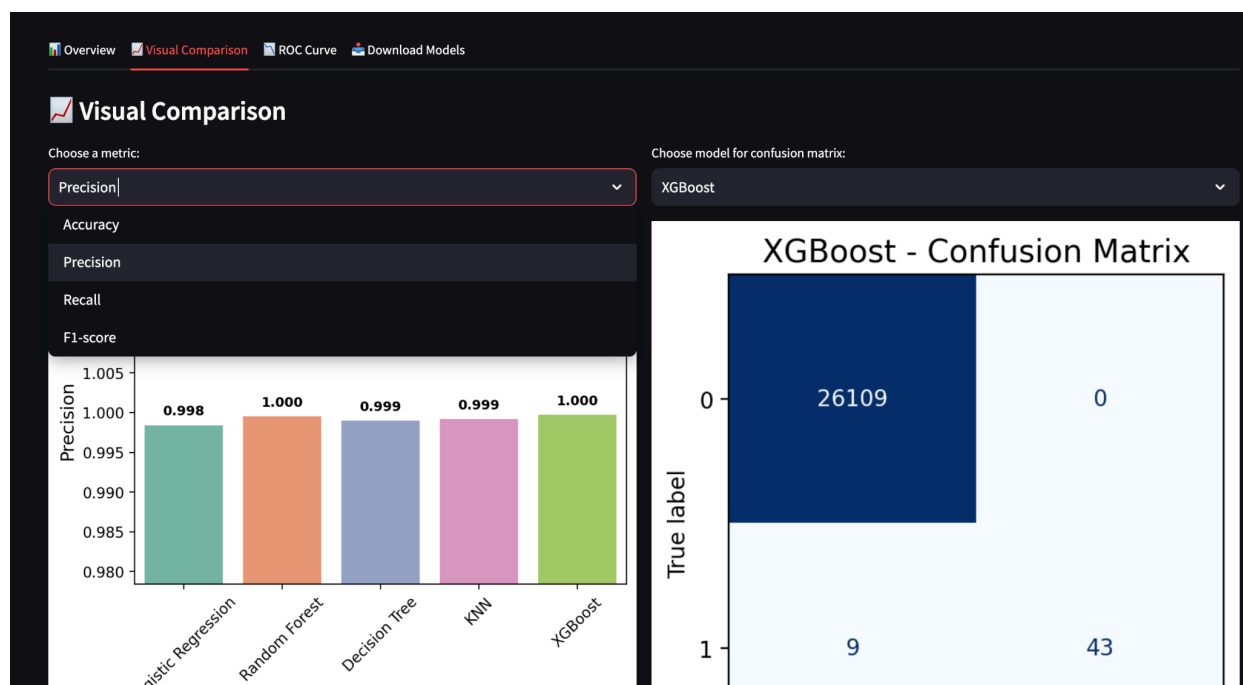
## **Dashboard (Streamlit) :**

We built a multi-tab Streamlit dashboard with the following features:

1. **Overview Tab:** Displays model evaluation metrics in a clean table
2. **Visual Comparison Tab:** Allows comparison of precision, recall, and F1-score across models using bar plots; also includes confusion matrix viewer
3. **ROC Curve Tab:** Lets users view and compare AUC-ROC curves model-wise
4. **Download Models Tab:** Provides trained `.pk1` files for download

 *Screenshots of the dashboard are provided in the next page.*





Deploy

## ML Model Comparison Dashboard

Overview Visual Comparison ROC Curve Download Models

### Download Trained Models

Click below to download any model you want to share or reuse.


Logistic Regression Random Forest SVM XGBoost

Made with ❤️ by Anushka, Manish & Jay

## ROC Curve (AUC)

Choose model for ROC curve:

Logistic Regression



## Key Features :



### Overview Tab

- Displays a table with metrics like **Precision**, **Recall**, and **F1-score** for each model.
- Styled for clarity and emphasis using font scaling and color.



### Visual Comparison Tab

- Users can select a metric (e.g., F1-score) from a dropdown.
- A bar plot shows how models compare.
- A separate dropdown loads the **Confusion Matrix** for any model.



### ROC Curve Tab

- Plots **ROC Curve** for a selected model.
- Displays **AUC score** visually for better model comparison.



### Download Tab

- Allows downloading of trained models in **.pk1** format.
- Great for portability or integration into other systems.

## Key Insights :

Through model experimentation, evaluation, and visualization, several key insights emerged:

### 1. XGBoost Outperformed Other Models

- After hyperparameter tuning, **XGBoost delivered the best balance between precision and recall.**
- It showed a strong **F1-score**, indicating reliability in detecting fraud while minimizing false alarms.

### 2. Precision vs Recall Trade-Off

- **Logistic Regression** and **SVM** had high precision but lower recall — they were more conservative in flagging fraud.
- **Random Forest** and **Gradient Boosting** provided better recall but risked more false positives.
- This highlighted the importance of **F1-score** and **ROC-AUC** as balanced evaluation metrics in imbalanced datasets.

### 3. Importance of Data Resampling (SMOTE)

- Using SMOTE drastically improved recall across all models, especially those that otherwise performed poorly on rare fraud examples.
- Without SMOTE, most models underfit the fraud class due to data imbalance.

### 4. Dashboard Helped Interpret Results Quickly

- The **bar charts, confusion matrices, and ROC curves** made it easy to compare models visually.
- Stakeholders or judges could quickly identify which model to trust and why.






## 5. Feature Importance Helped Understand Predictions

- SHAP values (used optionally during model tuning) revealed which features most influenced fraud detection.
- Although features were anonymized (V1–V28), their relative importance offered meaningful interpretability.

## Conclusion :

In this project, we successfully developed and evaluated a robust credit card fraud detection system using multiple machine learning models. We addressed the challenges of **class imbalance**, **model tuning**, and **evaluation interpretability** through a systematic pipeline.

Key achievements include:

-  **Training and comparing five classification models** on real-world fraud data
-  **Balancing the data using SMOTE** to handle extreme class imbalance
-  **Hyperparameter tuning** to optimize model performance
-  **Comprehensive evaluation** using F1-score, ROC-AUC, and confusion matrices
-  Building an **interactive Streamlit dashboard** to visualize and compare model performance

The final solution offers a **transparent, explainable, and easily deployable system** for fraud detection in financial datasets.

## Future Scope :

While the current solution is robust and production-ready, there are several opportunities to expand and improve this project:

### 1. Real-Time Prediction & Deployment

- Integrate the trained models into a real-time fraud detection pipeline using **APIs** or **batch prediction endpoints**.
- Serve the dashboard using **Streamlit Cloud**, **Heroku**, or **AWS** for public access.

### 2. Feature Engineering

- Investigate derived features like **transaction frequency**, **user velocity**, or **time of day**.
- These could enhance fraud detection accuracy with minimal overhead.

### 3. Advanced Algorithms

- Implement **LightGBM**, **CatBoost**, or **Neural Networks** to explore deeper learning techniques.
- Add **stacking** or **voting ensembles** to combine strengths of multiple models.

### 4. Precision-Recall Curves & Cost Analysis

- Add **PR curves** to complement ROC, especially in imbalanced settings.
- Conduct **cost-sensitive evaluation**, e.g., how much a false negative costs vs a false positive.



## 5. Data Privacy & Explainability

- Integrate **LIME/SHAP explanations** into the dashboard so users can interpret individual predictions.
- Consider techniques for **model auditing and bias analysis**, especially in financial applications.

## 6. User Upload Feature

- Let users upload new transaction CSVs in the dashboard and get predictions with fraud probability and SHAP explanations.

This roadmap would help evolve the current dashboard into a full-fledged ML-powered fraud detection platform.

**Video presentation link :**

[https://drive.google.com/file/d/1wevhGeFYOb8kEVRtqN1gAyoC0K6y1-ZU/view?usp=drive\\_link](https://drive.google.com/file/d/1wevhGeFYOb8kEVRtqN1gAyoC0K6y1-ZU/view?usp=drive_link)

